

Branch: master ▾

datasciencectacontent / markdown / rprog-breakingDownMakeVector.md

Find file

Copy path



lgreski minor tweaks

52cb0a3 on Aug 12, 2018

1 contributor

269 lines (158 sloc) 19.3 KB

Raw

Blame

History



Demystifying makeVector()

The second programming assignment in the Johns Hopkins University *R Programming* course on Coursera.org makes use of a prototype set of functions that illustrate caching of a mean from a vector. The overall objective of the assignment is to demonstrate the concept of lexical scoping. This assignment often confuses students because it is difficult for them to understand the concepts simply by looking at their implementation in code.

This article explains the code in the `cachemean.R` file, highlighting key R concepts and features that make the program work as expected. In addition to the concept of scoping, the assignment also introduces use of S3 objects without explicitly explaining how object orientation is implemented in R, causing large amounts of frustration in students when they are unable to get their implementations of `makeVector()` and `cacheSolve()` to work correctly.

What is lexical scoping?

Scoping is the mechanism within R that determines how R finds symbols (i.e. programming language elements) to retrieve their values during the execution of an R script.

R supports two types of scoping: lexical scoping and dynamic scoping. As noted by Hadley Wickham, "Dynamic scoping is primarily used within functions to save typing during interactive analysis," and will not be covered here. Hadley Wickham's [Advanced-R website's section on scoping issues](#) covers dynamic scoping.

Lexical scoping is used to retrieve values from objects based on the way functions are nested when they were written. Since *Programming Assignment 2* contains nested functions, to fully comprehend the assignment students must be able to visualize how the symbols are stored and accessed within the nested function. A more detailed explanation of lexical scoping is available on the [Advanced-R website's Function page](#), so I'll refer the reader there for the details.

Understanding of scoping is key to *R Programming Assignment 2*, because the fact that the "cache" works is due to how the code is built at design time, not how the code is called at runtime.

What's the importance of caching?

A cache is a way to [store objects in memory to accelerate subsequent access](#) to the same object. In statistics, some matrix algebra computations are notoriously expensive, such as calculating the inverse of a matrix. Therefore, if one needs to use the same inverted matrix for subsequent computations, it is advantageous to cache it in memory instead of repeatedly calculating the inverse. Programming Assignment 2 uses the scenario of needing to cache an inverted matrix as a way of illustrating how this might be done with a special matrix object that stores its inverse as an in-memory object.

The assignment uses an example of caching a mean to illustrate the underlying concepts of scoping and creating an S3 object, as I discuss in [R Objects, S Objects, and Lexical Scoping](#).

Overall Design of `makeVector()` and `cachemean()`

The `cachemean.R` file contains two functions, `makeVector()` and `cachemean()`. The first function in the file, `makeVector()` creates an R object that stores a vector and its mean. The second function, `cachemean()` requires an argument that is returned by `makeVector()` in order to retrieve the mean from the cached value that is stored in the `makeVector()` object's environment.

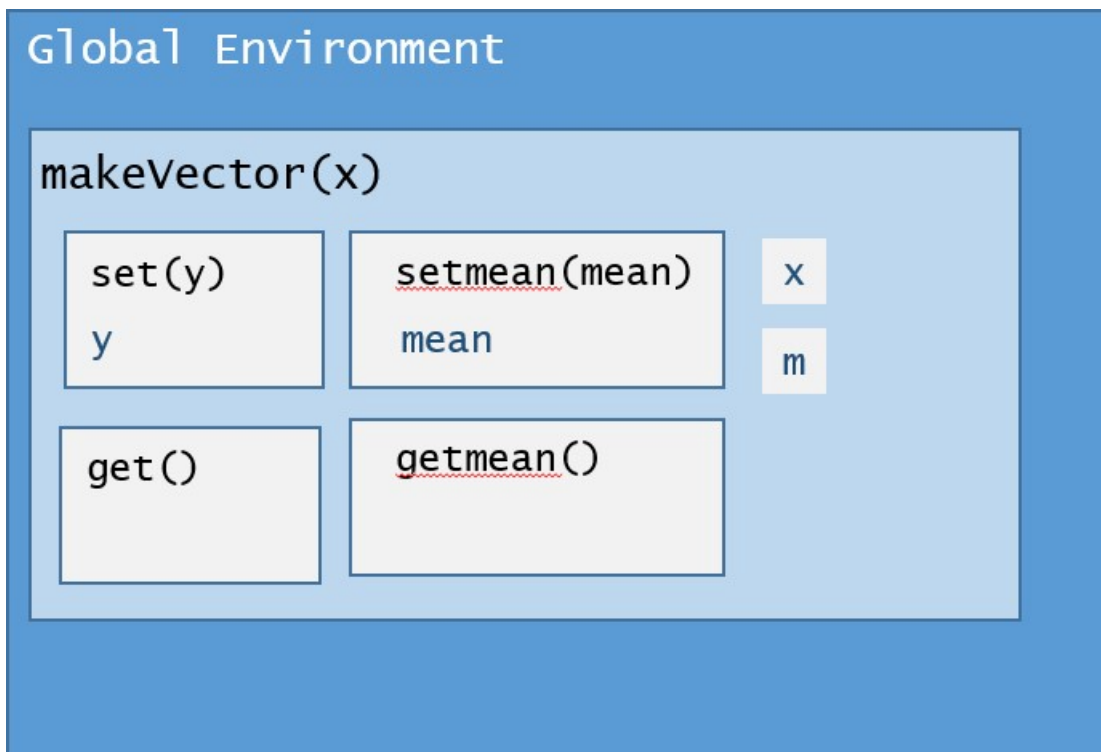
What's going on in `makeVector()`?

The key concept to understand in `makeVector()` is that it builds a set of functions and returns the functions within a list to the parent environment. That is,

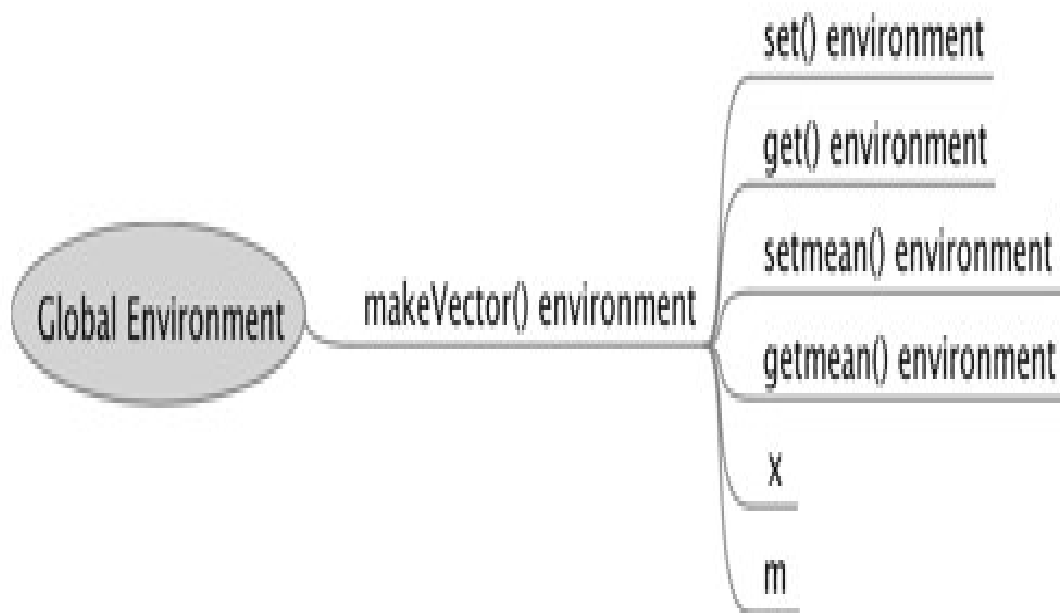
```
myVector <- makeVector(1:15)
```

results in an object, `myVector`, that contains four functions: `set()`, `get()`, `setmean()`, and `getmean()`. It also includes the two data objects, `x` and `m`.

Due to lexical scoping, `myVector` contains a complete copy of the environment for `makeVector()`, including any objects that are defined within `makeVector()` at design time (i.e., when it was coded). A diagram of the environment hierarchy makes it clear what is accessible within `myVector`.

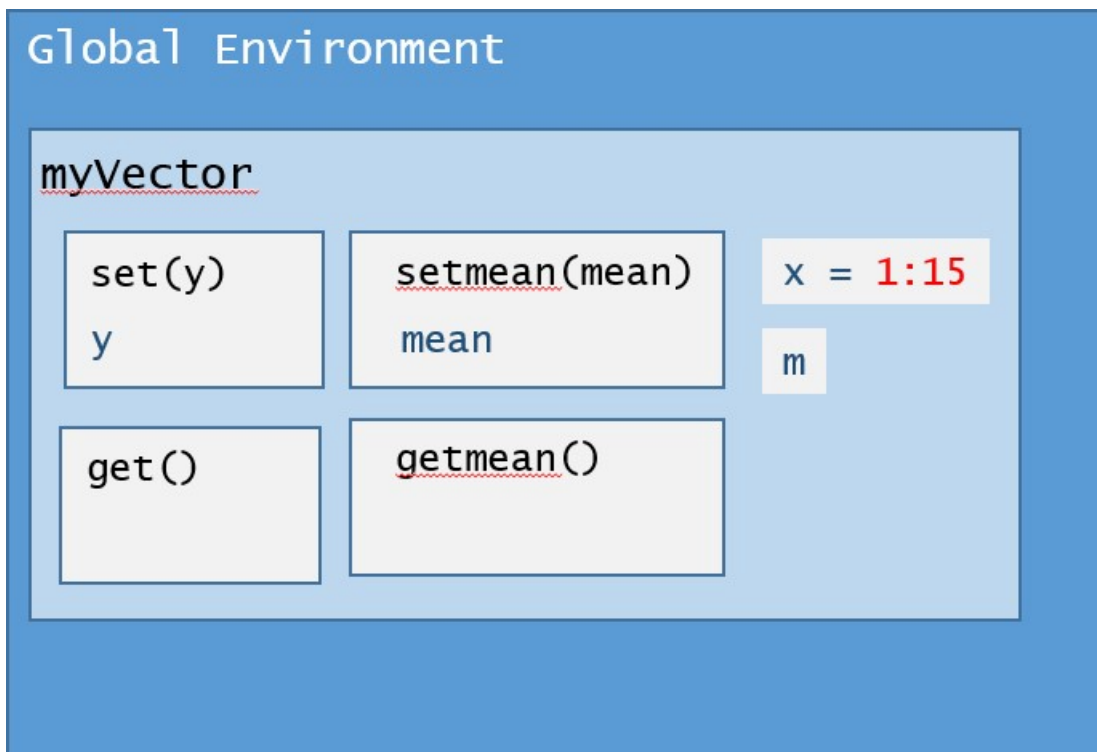


Illustrated as a hierarchy, the global environment contains the `makeVector()` environment. All other content is present in the `makeVector()` environment, as illustrated below.



Since each function has its own environment in R, the hierarchy illustrates that the objects `x` and `m` are siblings of the four functions, `get()`, `set()`, `getmean()`, and `setmean()`.

Once the function is run and an object of type `makeVector()` is instantiated (that is, created), the environment containing `myVector` looks like:



Notice that the object `x` contains the vector `1:15`, even though `myVector$set()` has not been executed. This is the case because the value `1:15` was passed as an argument into the `makeVector()` function. What explains this behavior?

When an R function returns an object that contains functions to its parent environment (as is the case with a call like `myVector <- makeVector(1:15)`), not only does `myVector` have access to the specific functions in its list, but it also retains access to the entire environment defined by `makeVector()`, including the original argument used to start the function.

Why is this the case? `myVector` contains pointers to functions that are within the `makeVector()` environment after the function ends, so these pointers prevent the memory consumed by `makeVector()` from being released by the garbage collector. Therefore, the entire `makeVector()` environment stays in memory, and `myVector` can access its functions as well as any data in that environment that is referenced in its functions.

This feature explains why `x` (the argument initialized on the original function call) is accessible by subsequent calls to functions on `myVector` such as `myVector$get()`, and it also explains why the code works without having to explicitly issue `myVector$set()` to set the value of `x`.

makeVector() step by step

Now, let's break the behavior of the function down, step by step.

Step 1: Initialize objects

The first thing that occurs in the function is the initialization of two objects, `x` and `m`.

```
makeVector(x = numeric()) {  
  m <- NULL  
  ...  
}
```

```
}
```

Notice that `x` is initialized as a function argument, so no further initialization is required within the function. `m` is set to `NULL`, initializing it as an object within the `makeVector()` environment to be used by later code in the function.

Furthermore, the `formals` part of the function declaration define the default value of `x` as an empty numeric vector. Initialization of the vector with a default value is important because without a default value, `data <- x$get()` generates the following error message.

```
Error in x$get() : argument "x" is missing, with no default
```

Step 2: Define the "behaviors" or functions for objects of type `makeVector()`

After initializing key objects that store key information within `makeVector()`, the code provides four basic behaviors that are typical for data elements within an [object-oriented program](#). They're called "getters and setters," and more formally known as [mutator and accessor](#) methods. As one might expect, "getters" are program modules that retrieve (access) data within an object, and "setters" are program modules that set (mutate) the data values within an object.

First `makeVector()` defines the `set()` function. Most of the "magic" in `makeVector()` takes place in the `set()` function.

```
set <- function(y) {
  x <- y
  m <- NULL
}
```

`set()` takes an argument that is named as `y`. It is assumed that this value is a numeric vector, but is not stated directly in the function `formals`. For the purposes of the `set()` function, it doesn't matter whether this argument is called `y`, `aVector` or any object name other than `x`. Why? Since there is an `x` object already defined in the `makeVector()` environment, using the same object name would make the code more difficult to understand.

Within `set()` we use the `<-` [form of the assignment operator](#), which assigns the value on the right side of the operator to an object in the parent environment named by the object on the left side of the operator.

When `set()` is executed, it does two things:

1. Assign the input argument to the `x` object in the parent environment, and
2. Assign the value of `NULL` to the `m` object in the parent environment. This line of code clears any value of `m` that had been cached by a prior execution of `cachemean()`.

Therefore, if there is already a valid mean cached in `m`, whenever `x` is reset, the value of `m` cached in the memory of the object is cleared, forcing subsequent calls to `cachemean()` to recalculate the mean rather than retrieving the wrong value from cache.

Notice that the two lines of code in `set()` do exactly the same thing as the first two lines in the main function: set the value of `x`, and `NULL` the value of `m`.

Second, `makeVector()` defines the getter for the vector `x`.

```
get <- function() x
```

Again, this function takes advantage of the lexical scoping features in R. Since the symbol `x` is not defined within `get()`, R retrieves it from the parent environment of `makeVector()`.

Third, `makeVector()` defines the setter for the mean `m`.

```
setmean <- function(mean) m <- mean
```

Since `m` is defined in the parent environment and we need to access it after `setmean()` completes, the code uses the `<-` form of the assignment operator to assign the input argument to the value of `m` in the parent environment.

Finally, `makeVector()` defines the getter for the mean `m`. Just like the getter for `x`, R takes advantage of lexical scoping to find the correct symbol `m` to retrieve its value.

```
getmean <- function() m
```

At this point we have getters and setters defined for both of the data objects within our `makeVector()` object.

Step 3: Create a new object by returning a list()

Here is the other part of the "magic" in the operations of the `makeVector()` function. The last section of code assigns each of these functions as an element within a `list()`, and returns it to the parent environment.

```
list(set = set, get = get,
     setmean = setmean,
     getmean = getmean)
```

When the function ends, it returns a fully formed object of type `makeVector()` to be used by downstream R code. One other important subtlety about this code is that each element in the list is [named](#). That is, each element in the list is created with a `elementName = value` syntax, as follows:

```
list(set = set,          # gives the name 'set' to the set() function defined above
     get = get,          # gives the name 'get' to the get() function defined above
     setmean = setmean,  # gives the name 'setmean' to the setmean() function defined
above
     getmean = getmean)  # gives the name 'getmean' to the getmean() function defined
above
```

Naming the list elements is what allows us to use the `$` [form of the extract operator](#) to access the functions by name rather than using the `[[` form of the extract operator, as in `myVector[[2]]()`, to get the contents of the vector.

Here it's important to note that the `cachemean()` function **REQUIRES** an input argument of type `makeVector()` . If one passes a regular vector to the function, as in

```
aResult <- cachemean(1:15)
```

the function call will fail with an error explaining that `cachemean()` was unable to access `$getmean()` on the input argument because `$` does not work with atomic vectors. This is accurate, because a primitive vector is not a list, nor does it contain a `$getmean()` function, as illustrated below.

```
> aVector <- 1:10
> cachemean(aVector)
Error in x$getmean : $ operator is invalid for atomic vectors
```

Explaining cachemean()

Without `cachemean()` , the `makeVector()` function is incomplete. Why? As designed, `cachemean()` is required to populate and/or retrieve the mean from an object of type `makeVector()` .

```
cachemean <- function(x, ...) {
  ...
```

Like `makeVector()` , `cachemean()` starts with a single argument, `x` , and an ellipsis that allows the caller to pass additional arguments into the function.

Next, the function attempts to retrieve a mean from the object passed in as the argument. First, it calls the `getmean()` function on the input object.

```
m <- x$getmean()
```

Then it checks to see whether the result is `NULL` . Since `makeVector()` sets the cached mean to `NULL` whenever a new vector is set into the object, if the value here is not equal to `NULL` , we have a valid, cached mean and can return it to the parent environment

```
if(!is.null(m)) {
  message("getting cached data")
  return(m)
}
```

If the result of `!is.null(m)` is `FALSE` , `cachemean()` gets the vector from the input object, calculates a `mean()` , uses the `setmean()` function on the input object to set the mean in the input object, and then returns the value of the mean to the parent environment by printing the mean object.

```
data <- x$get()
m <- mean(data, ...)
x$setmean(m)
m
```

Note that `cachemean()` is the only place where the `mean()` function is executed, which is why `makeVector()` is incomplete without `cachemean()`.

Putting the Pieces Together: How the functions work at runtime

Now that we've explained the design of each of these functions, here is an illustration of how they work when used in an R script.

```
aVector <- makeVector(1:10)
aVector$get()           # retrieve the value of x
aVector$getmean()       # retrieve the value of m, which should be NULL
aVector$set(30:50)      # reset value with a new vector
cachemean(aVector)      # notice mean calculated is mean of 30:50, not 1:10
aVector$getmean()       # retrieve it directly, now that it has been cached

> aVector <- makeVector(1:10)
> aVector$get() # retrieve the value of x
[1] 1 2 3 4 5 6 7 8 9 10
> aVector$getmean() # retrieve the value of m, which should be NULL
NULL
> # reset value with a new vector
> aVector$set(30:50)
> cachemean(aVector) # notice mean calculated is mean of 30:50, not 1:10
[1] 40
> aVector$getmean() # retrieve it directly, now that it has been cached
[1] 40
> |
```

Conclusion: what makes `cachemean()` work?

To summarize, the lexical scoping assignment in *R Programming* takes advantage of lexical scoping and the fact that functions that return objects of type `list()` also allow access to any other objects defined in the environment of the original function. In the specific instance of `makeVector()` this means that subsequent code can access the values of `x` or `m` through the use of getters and setters. This is how `cachemean()` is able to calculate and store the mean for the input argument if it is of type `makeVector()`. Because list elements in `makeVector()` are defined with names, we can access these functions with the `$` [form of the extract operator](#).

For additional commentary that explains how the assignment uses features of the S3 object system, please review [makeCacheMatrix\(\) as an Object](#).

Appendix A: What's the Point of this Assignment?

Once students get through the assignment, they frequently ask questions about its value and purpose. A good article explaining the value of lexical scoping in statistical computing is [Lexical Scoping and Statistical Computing](#), written by Robert Gentleman and Ross Ihaka at the University of Auckland.

Appendix B: `cachemean.R`

Here is the entire listing for `cachemean.R`.


```

makeVector <- function(x = numeric()) {
  m <- NULL
  set <- function(y) {
    x <- y
    m <- NULL
  }
  get <- function() x
  setmean <- function(mean) m <- mean
  getmean <- function() m
  list(set = set, get = get,
       setmean = setmean,
       getmean = getmean)
}
cachemean <- function(x, ...) {
  m <- x$getmean()
  if(!is.null(m)) {
    message("getting cached data")
    return(m)
  }
  data <- x$get()
  m <- mean(data, ...)
  x$setmean(m)
  m
}

```

Appendix C: Frequently Asked Questions

Q: Why doesn't `cachemean()` return the cached value? My code looks like:

```

cachemean(makeVector(1:100))
cachemean(makeVector(1:100))

```

A: Code written this way creates two different objects of type `makeVector()`, so the two calls to `cachemean()` initialize the means of each instance, rather than caching and retrieving from a single instance. Another way of illustrating how the above code operates is as follows.

```

> source('C:/Users/leona/gitrepos/datascience/rprogramming/cachemean.R')
> #
> # cachemean examples
> #
> a <- makeVector(1:100)
> b <- makeVector(1:100)
>
> # now call cachemean twice to show retrieval from cache
> cachemean(a)
[1] 50.5
> cachemean(a)
getting cached data
[1] 50.5
>
> # now call cachemean twice to show retrieval from cache
> cachemean(b)
[1] 50.5
> cachemean(b)
getting cached data
[1] 50.5
> |

```

Notice how the first call to `cachemean()` sets the cache, and the second call retrieves data from it.

Q: Why is `set()` never used in the code?

A: `set()` is included so that once an object of type `makeVector()` is created, its value can be changed without initializing another instance of the object. It is unnecessary the first time an object of type `makeVector()` is instantiated. Why? First, the value of `x` is set as a function argument, as in `makeVector(1:30)`. Then, the first line of code in the function sets `m <- NULL`, simultaneously allocating memory for `m` and setting it to `NULL`. When a reference to this object is passed to the parent environment when the function ends, both `x` and `m` are available to be accessed by their respective `get` and `set` functions.

The following code illustrates the use of `set()`.

```
> a <- makeVector(1:10)
>
> # now call cachemean twice to show retrieval from cache
> cachemean(a)
[1] 5.5
> cachemean(a)
getting cached data
[1] 5.5
>
>
> a$set(c(30,50,90,120,180))
>
> # now call cachemean twice to show retrieval from cache
> cachemean(a)
[1] 94
> cachemean(a)
getting cached data
[1] 94
> |
```

Q: Why is `x` set with a default value in `makeVector()` ?

A: Since `x` is an argument, the only place where one can set a default for it is in the formals. The type of error returned by `cachemean()` when a default value is not set,

```
Error in x$get() : argument "x" is missing, with no default
```

is undesirable. Our code should directly handle error conditions rather than relying on the underlying error handling in R.

It's perfectly valid to create an object of type `makeVector()` without populating its value during initialization. `makeVector()` includes a setter function so one can set its value after the object is created. However, the object must have valid data, a numeric vector, prior to executing `cachemean()`.

Ideally, `cachemean()` would include logic to validate that `x` is not empty prior to calculating a mean. The default setting of `x` enables `cachemean()` to return `NaN`, which is a reasonable result.

References

1. Chi, Yau -- [R-Tutor Named List Members](#), retrieved July 20, 2016.
2. Wickham, Hadley -- [Advanced-R Functions](#), retrieved July 17, 2016.
3. Wickham, Hadley -- [Advanced-R Scoping Issues](#), retrieved July 17, 2016.